

Query Processing



Steps of Query Processing

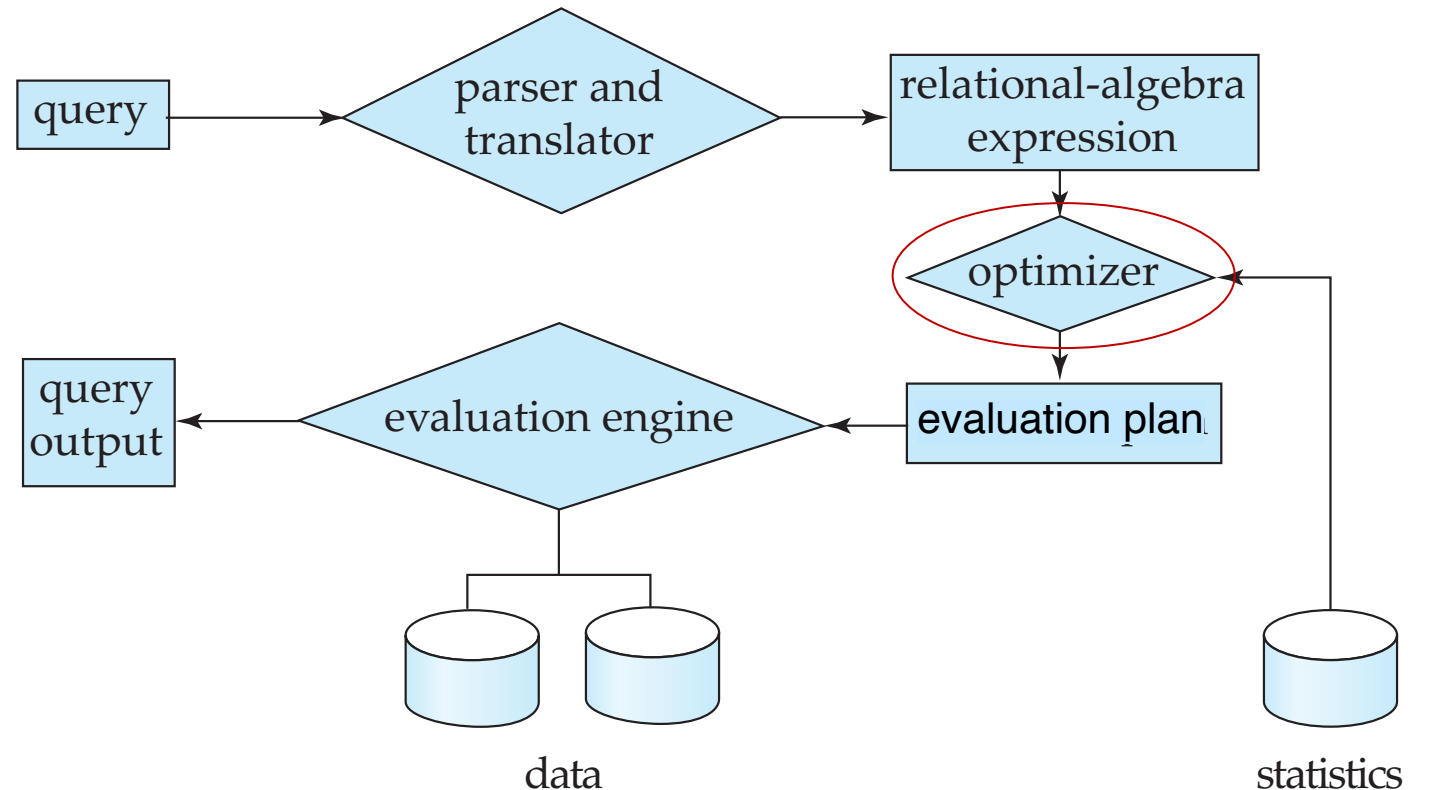
1. Parsing and translation

- Parser to check syntax, verify relation names, etc.
- Translate query into internal form (~ relational algebra)

2. Optimization

1. Equivalent expression trees
2. Best evaluation plan

3. Evaluation



Outline


- Query Optimization
- Generation Equivalent Expressions
 - Equivalence Rules
- Cost Estimation of Expressions

Query optimization

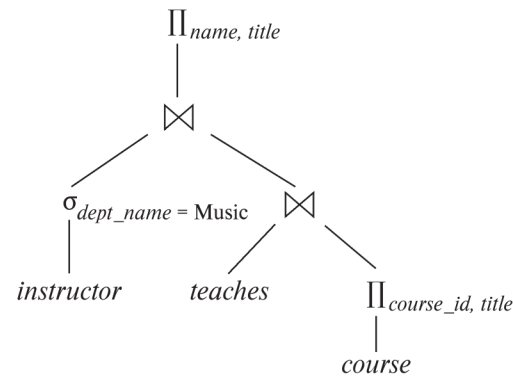
- Single relational algebra expression ~ many equivalent expressions
 - Equivalent expressions trees might involve different computational cost

SELECT salary FROM instructor WHERE salary < 75000

$\sigma_{\text{salary} < 75000}(\Pi_{\text{salary}}(\text{instructor}))$
 $\Pi_{\text{salary}}(\sigma_{\text{salary} < 75000}(\text{instructor}))$




(a) Initial expression tree




(b) Transformed expression tree

Query optimization

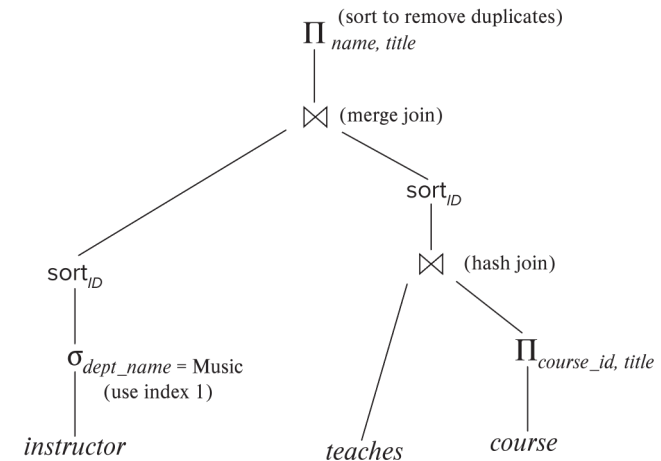
- Different algorithms can be used to evaluate the same relational algebra operation
 - **Evaluation primitive**: A relational algebra operation annotated with instructions on how to evaluate it
 - **Evaluation plan**: Sequence of primitives to evaluate a query
 - Feed to the *query-evaluation engine*



(a) Initial expression tree



(b) Transformed expression tree



(c) Query evaluation plan

Query optimization

- **Query optimization**: To select the most efficient **evaluation plan** among all the possible ones to process a query
 - User is not expected to write the most efficient query!
- Worth taking some time to select a good plan!
 - Cost difference may be huge... From seconds to days!
- Steps for query optimization:
 1. Generate **equivalent expression trees** using **equivalence rules**
 - Expected to find (better) expressions at relational-algebra level
 2. Annotate resultant expressions to get alternative evaluation plans
 - Detailed strategy for processing the query (algorithms, indices, etc.)
 3. Choose the cheapest **evaluation plan** based on **estimated costs**



Query optimization

- Do you want to see the actual query evaluation plan?
 - Use **explain** *<query>* to display plan chosen by query optimizer, along with cost estimates
Syntax in your DBMS may change!
 - Use **explain analyse** *<query>* to see actual runtime statistics found by running the query

Generating Equivalent Expressions

- Two relational algebra expressions are **equivalent** if both generate the same set of tuples on every *legal* DB instance
 - The order is irrelevant
 - SQL considers *multisets* of tuples
 - Two expressions are *equivalent* if both generate the same *multiset* of tuples on every *legal* DB instance.
- The *optimizer* uses equivalence rules to obtain equivalent expressions
 - An **equivalence rule** states that expressions of two forms are equivalent

Equivalence Rules: some examples

1. **Conjunctive selection** is similar to *sequence of selections*:

$$\sigma_{\theta_1 \wedge \theta_2}(E) \equiv \sigma_{\theta_1}(\sigma_{\theta_2}(E))$$

2. **Selection** operation is *commutative*:

$$\sigma_{\theta_1}(\sigma_{\theta_2}(E)) \equiv \sigma_{\theta_2}(\sigma_{\theta_1}(E))$$

3. In a **sequence of projections**, only the last one matters:

$$\Pi_{L_1}(\Pi_{L_2}(\dots(\Pi_{L_n}(E))\dots)) \equiv \Pi_{L_1}(E)$$

where $L_1 \subseteq L_2 \dots \subseteq L_n$

4. **Selections** can be combined into **Cartesian products and theta joins**.

$$\begin{aligned}\sigma_{\theta}(E_1 \times E_2) &\equiv E_1 \bowtie_{\theta} E_2 \\ \sigma_{\theta_1}(E_1 \bowtie_{\theta_2} E_2) &\equiv E_1 \bowtie_{\theta_1 \wedge \theta_2} E_2\end{aligned}$$


Equivalence Rules: some examples

5. Theta (and natural) join is *commutative*:

$$E_1 \bowtie E_2 \equiv E_2 \bowtie E_1$$

6. a) Natural join is *associative*:

$$(E_1 \bowtie E_2) \bowtie E_3 \equiv E_1 \bowtie (E_2 \bowtie E_3)$$



b) Theta join is *associative* only when:

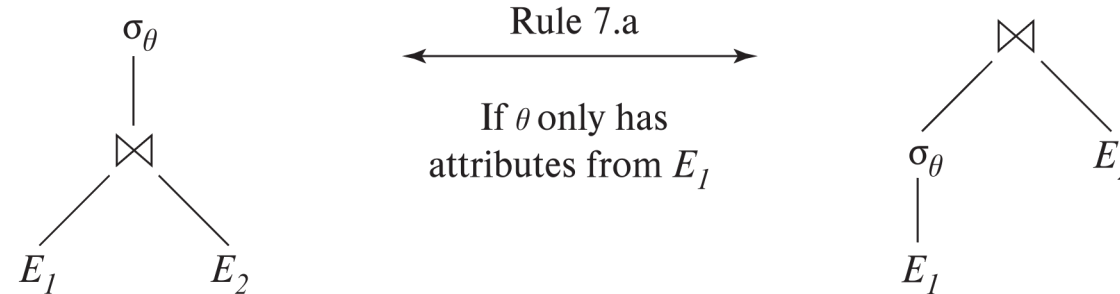
$$(E_1 \bowtie_{\theta_1} E_2) \bowtie_{\theta_2 \wedge \theta_3} E_3 \equiv E_1 \bowtie_{\theta_1 \wedge \theta_3} (E_2 \bowtie_{\theta_2} E_3)$$

where θ_2 involves attributes from only E_2 and E_3 .

Equivalence Rules: some examples

7. The **selection** operation *distributes* over the **theta join** when all the attributes in θ_s come from a single join expression:

$$\sigma_{\theta_s}(E_1 \bowtie_{\theta_j} E_2) \equiv (\sigma_{\theta_s}(E_1)) \bowtie_{\theta_j} E_2$$



8. The **projection** operation *distributes* over the **theta join** if the join condition θ involves only attributes from $L_1 \cup L_2$:

$$\Pi_{L_1 \cup L_2}(E_1 \bowtie_\theta E_2) \equiv \Pi_{L_1}(E_1) \bowtie_\theta \Pi_{L_2}(E_2)$$

where L_1 and L_2 are attributes from E_1 and E_2 , respectively.

And many more...

Generating Equivalent Expressions

- How is all this useful?
- Consider **join associativity**:

$$(r_1 \bowtie r_2) \bowtie r_3 \equiv r_1 \bowtie (r_2 \bowtie r_3)$$

- We can decide to compute **first** joins that lead to **smaller intermediate** relations

E.g.,

$$\Pi_{name, title} \left(\sigma_{dept_name = \text{“Music”}} (instructor) \bowtie teaches \bowtie \Pi_{course_id, title} (course) \right)$$

Only a few instructors belong to *Music* dept., so start with

$$(\sigma_{dept_name = \text{“Music”}} (instructor) \bowtie teaches)$$



Generating Equivalent Expressions

E.g., “Find all instructors’ names of Music dept. who taught a course in 2017, and the title of these courses”

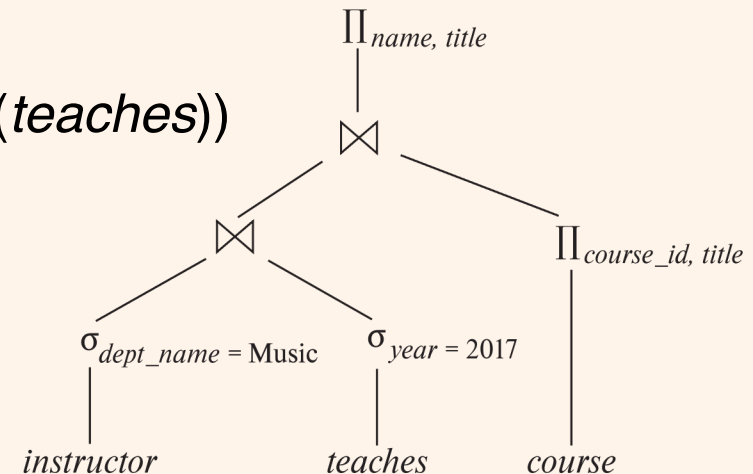
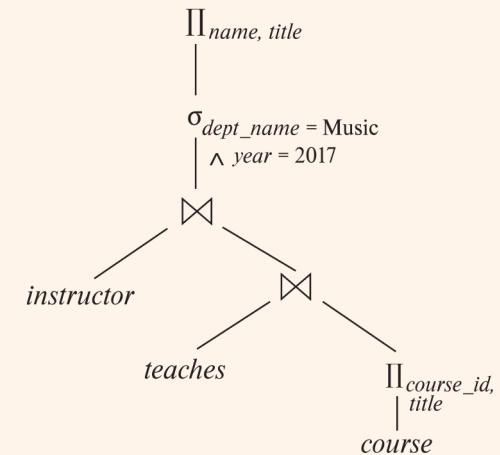
$$\Pi_{name, title}(\sigma_{dept_name = \text{“Music”} \wedge year = 2017} (instructor \bowtie (teaches \bowtie \Pi_{course_id, title}(course))))$$

- By **join associativity** (ER 6a):

$$\Pi_{name, title}(\sigma_{dept_name = \text{“Music”} \wedge year = 2017} (instructor \bowtie teaches) \bowtie \Pi_{course_id, title}(course))$$

- By “**conjunctive selection to sequence of selections**” (ER 1) and “**selection distribution on join**” (ER 7):

$$\Pi_{name, title}(\sigma_{dept_name = \text{“Music”}}(instructor) \bowtie \sigma_{year = 2017}(teaches) \bowtie \Pi_{course_id, title}(course))$$



Generating Equivalent Expressions

- Query optimizers perform systematic **enumeration of equivalent expressions**

- Simplest approach (**expensive!!**):

```
genAllEquivExpressions(E) :
```

```
1.  $T \leftarrow \{E\}$ 
```

```
2. while T grows then
```

```
    1. Check equivalence rules  $R_j$  in expressions  $E_i$  in T
```

```
    2. if any subexpression  $e_i$  of  $E_i$  matches one side of  $R_j$  then
```

```
        1.  $E' = E_i - R_j^1 + R_j^2$            //  $E_i$  but substituting  $R_j$  sides
```

```
        2. if  $E'$  not in T then     $T \leftarrow T \cup \{E'\}$ 
```

```
3. return T
```

- Two improvements:

- Use shared subexpressions to reduce space requirements
- Use **cost estimates** to avoid certain expressions

Now what? Or which?

- Now we know how to **generate** tons of equivalent **expressions**...

What's next?

- Can we tell **which expression is better**?

Let's estimate their cost

Less cost = more efficient query processing




Cost Estimation of Expressions

- **Cost** of an expression depends on the **size** of the input relation(s) and **other statistics**.
 - We use **approximated** values, and so is the cost
 - Recomputed after each DB modification (**expensive!!**)
 - Recomputed during idle time on specific events or under demand (**outdated values used!!**)
- In an expression tree, estimation of relevant statistics starts **from the leaves up to the root**.

Statistics:


- n_r : no. tuples in relation r
- l_r : size of tuples of relation r
- $v_{r(A)}$: no. distinct values of the (set of) attribute(s) A in relation r
- b_r : no. blocks that contain tuples of relation r
- f_r : blocking factor of relation r (no. tuples of r that fit into a block)
 - If tuples are physically stored together: $b_r = \lceil n_r / f_r \rceil$
- Info about indices (e.g., tree depth, no. leaves)



Cost Estimation of Expressions

- **Histograms** are usually stored too on numeric *attributes*

- Equal-width vs Equal-depth
- Frequency
- No. distinct values
- Without no. distinct values, cost estimates assume uniform distribution



- Many DB systems store the **N most-frequent values** with their counts
 - Histogram is built only on remaining values

Size Estimation of *Selection*

- A **single equality comparison**: $\sigma_{A=a}(r)$
 - $c = n_r / v_{r(A)}$: no. records that would satisfy the selection
No information \rightarrow uniformity
 - $c = 1$: if A is a key attribute
 - $c = f_{h_a} / v_{h_a}$: Use histogram when available
- A **single comparison**: $\sigma_{A \leq a}(r)$ [$\sigma_{A \geq a}(r)$ is symmetric]
 - If $\min_{r(A)}$ and $\max_{r(A)}$ are stored in our statistic's catalog:
 - $c = 0$, if $a < \min_{r(A)}$
 - $c = n_r$, if $a \geq \max_{r(A)}$
 - $c = n_r \cdot (a - \min_{r(A)}) / (\max_{r(A)} - \min_{r(A)})$, otherwise
 - Use histograms when available
 - $c = n_r / 2$: if value a is unknown or in absence of statistical information

Size Estimation of *Join*

- **Cartesian product** ($r \times s$):
 - $n_x = n_r \cdot n_s$ tuples, where each tuple takes $l_x = l_r + l_s$ bytes
- **Natural join** ($r \bowtie s$): Let us assume $r(R)$ and $s(S)$
 - If $R \cap S = \emptyset$, then $(r \bowtie s) \equiv (r \times s)$
 - If $R \cap S$ is a foreign key in S referencing R , then
no. tuples in $(r \bowtie s) =$ no. tuples in s
 - If $R \cap S = \{A\}$ is not a key in any case, assuming *value uniformity*:
$$n_{\bowtie} = \min(n_r \cdot n_s / v_{r(A)} ; n_r \cdot n_s / v_{s(A)})$$
- **Theta join** ($r \bowtie_{\theta} s$): calculate it from $(r \bowtie_{\theta} s) \equiv \sigma_{\theta}(r \times s)$

With histograms,
these estimates
can be improved

Number of Distinct Values Estimation

■ Selection: $\sigma_{\theta}(r)$

- If θ forces A to take a specific value (e.g., $\theta \equiv (A = 3)$):

$$V_{\sigma_{\theta}(r)(A)} = 1$$

- If θ forces A to take a value on a set (e.g., $\theta \equiv (A = 1) \vee (A = 3)$):

$$V_{\sigma_{\theta}(r)(A)} = \text{size of the given value set}$$

- If $\theta \equiv (A \text{ op } r)$:

*** op = comparison operator*

$$V_{\sigma_{\theta}(r)(A)} = V_{r(A)} \cdot S$$

- In general: $V_{\sigma_{\theta}(r)(A)} = \min(V_{r(A)}, n_{\sigma_{\theta}(r)})$

■ Join: $r \bowtie s$

- If all attributes in A are from r :

$$V_{(r \bowtie s)(A)} = \min(V_{r(A)} ; n_{(r \bowtie s)})$$

- If A contains attributes A_r and A_s from r and s :

$$V_{(r \bowtie s)(A)} = \min(V_{r(A_r)} \cdot V_{s(A_s - A_r)} ; V_{r(A_r - A_s)} \cdot V_{s(A_s)} ; n_{r \bowtie s})$$

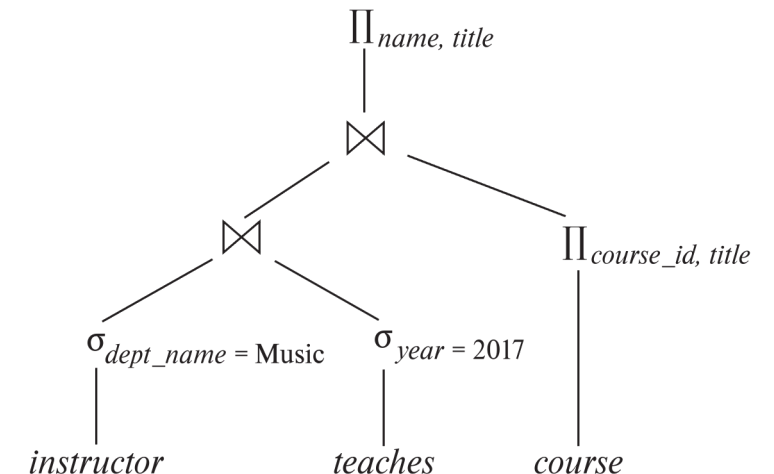
$A \equiv \text{attribute}(s)$
obtained after the join

Cost Estimation of Expressions

- **Cost** of an expression depends on the **size** of the input relation(s) and **other statistics**.
 - We use **approximated** values, and so is the cost
 - Recomputed after each DB modification (**expensive!!**)
 - Recomputed during idle time on specific events or under demand (**outdated values used!!**)
- In an expression tree, estimation of relevant statistics starts **from the leaves up to the root**.

Statistics:

- n_r : no. tuples in relation r
- l_r : size of tuples of relation r
- $v_{r(A)}$: no. distinct values of the (set of) attribute(s) A in relation r
- b_r : no. blocks that contain tuples of relation r
- f_r : blocking factor of relation r (no. tuples of r that fit into a block)
 - If tuples are physically stored together: $b_r = \lceil n_r / f_r \rceil$
- Info about indices (e.g., tree depth, no. leaves)



Query Processing

