

Transactions



Outline

- Transactions
- Schedules
- Properties
 - Serializability
 - Recoverability
 - Cascadeless
- Isolation Levels

Transactions

- **Transaction**: set of operations (access/update of diff. data items) that forms a *single logical unit of work*.

- E.g., transfer \$50 from account A to account B:

1.	read (A)	4.	read (B)
2.	$A := A - 50$	5.	$B := B + 50$
3.	write (A)	6.	write (B)

- Definition:

```
begin transaction  
...SQL statements...  
end transaction
```

Transactions

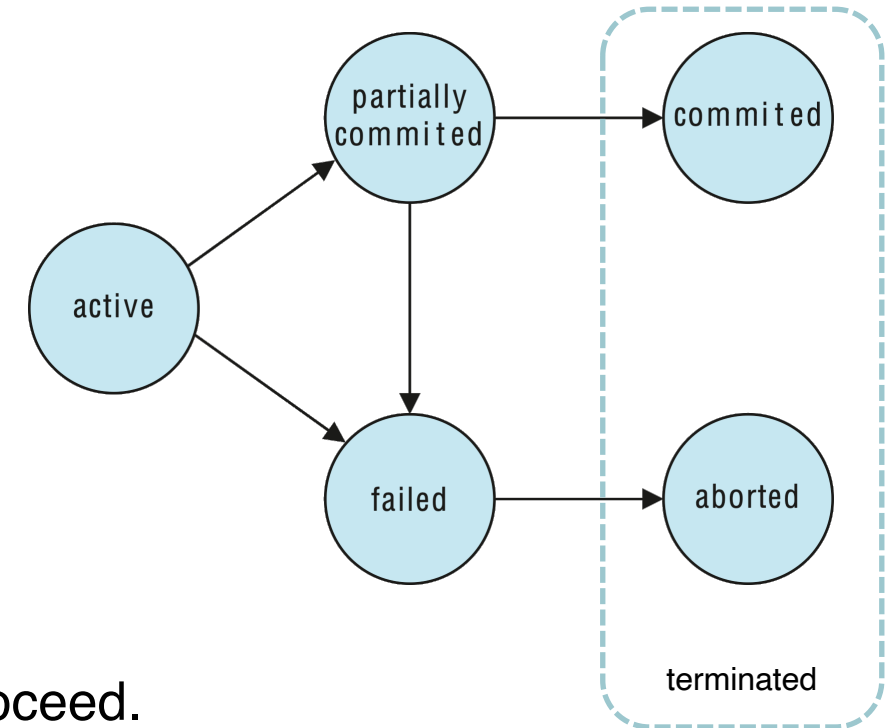
- Transactions go through different *states*:

- **Active** (initial state): a transaction is in this state during execution.
- **Partially committed**: after the transaction's final operation has been executed.
- **Committed**: after successful completion.

The DB is now in a **new** persistent and consistent state

- **Failed**: once discovered that execution cannot proceed.
 - Any change made must be *undone* (**roll back** transaction)
How? E.g., maintain a *log* with all info. needed for rolling back
- **Aborted**: transaction didn't complete its execution successfully.
2 possibilities: restart or definitely kill the transaction

- **Terminated**: Committed or aborted



Transactions

- **Transaction**: set of operations (access/update of diff. data items) that forms a *single logical unit of work*.
 - E.g., transfer \$50 from account A to account B:

1. <code>read(A)</code>	4. <code>read(B)</code>
2. <code>A := A - 50</code>	5. <code>B := B + 50</code>
3. <code>write(A)</code>	6. <code>write(B)</code>
- *Potential issues*:
 - *Failures* (e.g., hardware failures, system crashes, ...)
 - *Concurrent* execution of multiple transactions
- **Transaction manager** allows us to focus on transaction definition, ignoring these issues.



Transactions

- **Transaction**: set of operations (access/update of diff. data items) that forms a *single logical unit of work*.
 - E.g., transfer \$50 from account A to account B:

1. <code>read(A)</code>	4. <code>read(B)</code>
2. <code>A := A - 50</code>	5. <code>B := B + 50</code>
3. <code>write(A)</code>	6. <code>write(B)</code>
- Properties (*ACID*):
 - **Atomicity** [all-or-none]: *all* the operations are *fully executed*, or, in case of *failure*, partial results must be *undone*
 - **Consistency**: explicit/implicit constraints are fulfilled
Total money in the bank is the same before and after transaction
 - **Isolation**: a transaction shouldn't see partially modified data of a *concurrent* transaction, only after completion
 - **Durability**: after transaction is completed, its effects must persist (in stable storage).
Don't want to lose the transaction's result within a failure.



Transactions' Schedules

- To ensure *consistency*, it is far easier to run transactions **serially**
- Why then transactions are run **concurrently**?
 - To improve *performance* and resource *utilization*
E.g., a transaction can use the CPU while another one reads/writes to disk
 - To *reduce average response time*
E.g., *short* transactions don't wait for previous *long* ones to complete.
- **Schedule**: chronological ordering to execute (interleaved) operations from different concurrent transactions
 - Include all instructions from all involved transactions
 - Preserve order of operations in each individual transaction

Concurrency-control schemes
create schedules
which preserve
consistency and
isolation

Next lesson



Transactions' Schedules

- E.g., T_1 is a transfer of \$50 from A to B ; and T_2 a transfer of 10% of the balance from A to B .

Serial schedules

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

$A = 855$

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

$A = 850$



Transactions' Schedules

- E.g., T_1 is a transfer of \$50 from A to B ; and T_2 a transfer of 10% of the balance from A to B .

Concurrent schedules

T_1	T_2
read (A) $A := A - 50$ write (A)	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	read (B) $B := B + temp$ write (B) commit

The sum $A + B$ is preserved

Consistent!

T_1	T_2
read (A) $A := A - 50$	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B)
write (A) read (B) $B := B + 50$ write (B) commit	read (B) $B := B + temp$ write (B) commit

The sum $A + B$ is not preserved

Inconsistent!



Property I: Serializability

- **Basic assumption:**

Each transaction **individually** preserves DB consistency



Serial schedules preserve DB consistency

- A *concurrent schedule* is **serializable** if it is *equivalent* to a serial schedule.

- For scheduling, only **read(Q)** and **write(Q)** are significant operations

- How to check if two schedules are *equivalent*??

- **Conflict serializability**
- View serializability



Property I: Conflict Serializability

- Consider instructions I_i, I_j from transactions T_i, T_j in the same schedule:
 - I_i and I_j are **exchangeable** if they refer to *different* data points, $I_i(\underline{P})$ and $I_j(\underline{Q})$
 - I_i and I_j **conflict** if they work on the *same* data point Q and at least one is a **write**:
 - If $I_i = \text{read}(Q), I_j = \text{write}(Q)$, the order of I_i and I_j **matters**
 - If $I_i = \text{write}(Q), I_j = \text{read}(Q)$, order **matters**
 - If $I_i = \text{write}(Q), I_j = \text{write}(Q)$, order **matters**
 - If $I_i = \text{read}(Q), I_j = \text{read}(Q)$, order **doesn't matter**

T_1	T_2
read(A)	
write(A)	
read(B)	
write(B)	
	read(A)
	write(A)
	read(B)
	write(B)

T_1	T_2
read(A)	
write(A)	
	read(A)
	write(A)
read(B)	
write(B)	
	read(B)
	write(B)

T_1	T_2
read(A)	
write(A)	
	read(A)
read(B)	
	write(A)
write(B)	
	read(B)
	write(B)

Property I: Conflict Serializability

- **Conflicts** impose an *order* among instructions
 - If I_i and I_j are **non-conflicting instructions**, schedules $S\langle I_i, I_j \rangle$ and $S'\langle I_j, I_i \rangle$ return the same result.
- Schedule S is:
 - **conflict equivalent** to S' if we can transform S into S' by a series of swaps of *non-conflicting instructions*
 - **conflict serializable** if it is *conflict equivalent* to a *serial schedule* S^s


T_1	T_2	T_1	T_2	T_1	T_2
read(A)		read(A)		read(A)	
write(A)		write(A)		write(A)	
read(B)			read(A)	read(B)	read(A)
write(B)			write(A)		write(A)
	read(A)	read(B)		write(B)	
	write(A)	write(B)			read(B)
	read(B)		read(B)		write(B)
	write(B)		write(B)		

Different *serial schedules* of a transaction set are not always conflict equiv.


Property I: Conflict Serializability

- **Testing** serializability of a schedule S for transactions T_1, T_2, \dots, T_n
- **Precedence graph**: *directed* graph where vertices are transactions and an arc $T_i \rightarrow T_j$ sets an order (T_i comes before T_j)
 - Given $T_i \rightarrow T_j$, any serial schedule S^s equivalent to S runs T_i before T_j .

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit




T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit





T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Inconsistent !!



Property I: Conflict Serializability

- **Testing** serializability of a schedule S with a *precedence graph*, G , is simple:
 - schedule S is *conflict serializable* if G is *acyclic*
 - Classical algorithms that look for **cycles** in a graph G run in $O(n^2)$ time
with n = no. transactions



Exercise: check conflict serializability

- Is this schedule conflict serializable?

T ₁	T ₂
read(A) read(B) write(B) write(A)	 read(A) write(A)

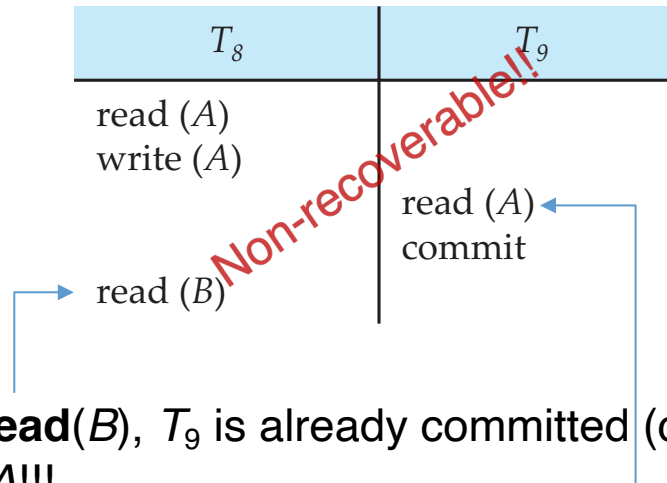
T ₃	T ₄
read(A) read(B) write(B) write(A)	 read(A) write(C)

T ₃	T ₄
read(A) read(B) write(B) write(A)	 read(B) write(B)

T ₃	T ₄
read(C) read(B) write(B) write(A)	 read(A) write(C)

Property II: Recoverable Schedules

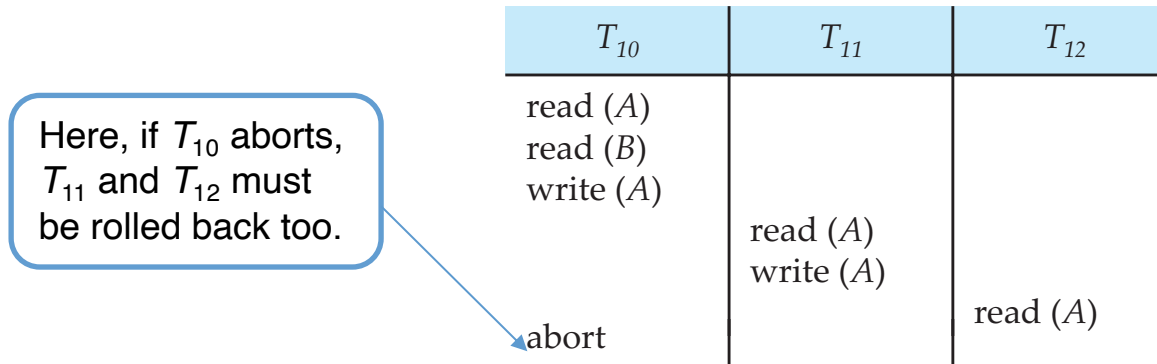
- What happens if a concurrent transaction fails?
 - Schedules must be **recoverable**
- **Recoverable schedule**: All the transactions are recoverable.
 - If transaction T_b *reads* data previously *written* by transaction T_a , the **commit** of T_a appears before that of T_b .



- Here, if T_8 fails after **read(B)**, T_9 is already committed (cannot be aborted), and it used an **inconsistent** value of A!!!

Property III: Cascadeless Schedules

- **Cascading rollback**: In recoverable schedules, a *single transaction failure* can lead to the rollback of a set of dependent transactions
 - Can lead to undo a significant amount of work



- In **cascadeless schedules**, cascading rollbacks cannot occur
 - If transaction T_b reads data previously written by T_a , the commit of T_a appears before the **read** of T_b .

Cascadeless is a **stronger** condition than *recoverable*

Summary

- Serializability:
 - A schedule of concurrent transactions is equivalent to a serial schedule

- Recoverability:
 - The transactions of a schedule are all recoverable; that is, a previous consistent state can be reached by rolling back

- Cascadeless:
 - As recoverability, but the previous consistent state is reached by rolling back **only a minimum** no. concurrent transactions

Transaction Isolation Levels

- **Serializability** ensures that concurrent transactions maintain consistency
 - To ensure **serializability**, too **little concurrency** can be allowed
 - If possible inconsistencies are not relevant for the app, **weaker levels of consistency** might be acceptable

E.g., long transactions whose result doesn't need to be precise, such as DB statistics computed for query optimization
- **Isolation levels:**
 - **Serializable**: usually ensures serializable execution.
 - **Repeatable read**: only reads committed data, and two reads of a data item by a transaction returns the same result.

Phantom reads are possible.
 - **Read committed**: allows only committed data to be read, but does not ensure repeatable reads.

I.e., two reads of a data item might be different.
 - **Read uncommitted**: allows uncommitted data to be read.

Lowest isolation level allowed by SQL

No isolation level allows **dirty writes**



Transaction Isolation Levels

- *Serializability* ensures that concurrent transactions maintain consistency

- To ensure **serializability**, too **little concurrency** can be allowed

- If possible inconsistencies are not relevant for the app **weaker levels of**

Trade-off: isolation vs. concurrency

You might accept a weaker isolation level to improve DBMS performance.

- **Isolation levels:**

- **Serializable:** Usually implemented by RDBMS

- **Repeatable:** + and two reads of a data item by a

transaction re **auto-commit**

Phantom reads are possible

- **Read committed:** allows only committed data to be read, but does not ensure repeatable reads.

I.e., two reads of a data item might be different.

- **Read uncommitted:** allows uncommitted data to be read.

Lowest isolation level allowed by SQL

No isolation level allows **dirty writes**

SQL Statements as Transactions

- E.g.,

Transaction 1 \equiv **select** *ID, name* **from** *instructor* **where** *salary* > 1900

Transaction 2 \equiv **insert into** *instructor* **values** ('11', Joe', 'Marketing', 1970)

Do T_1 and T_3 conflict?

- Result of T_1 is different depending on the order (T_1, T_2) or (T_2, T_1)
- In case (T_2, T_1), there is a clear precedence: $T_2 \rightarrow T_1$
- In case (T_1, T_2), a data point for new instructor ($ID=11$) doesn't even exist!

- **Phantom phenomenon**

- Consider now (Wu's salary = 1890):

Transaction 3 \equiv **update** *instructor* **set** *salary=salary*1.1* **where** *name='Wu'*

Do T_1 and T_3 conflict?

- Consider conflicts between *insert/delete/update* statements and *select*'s where-**predicates**, (e.g., "*salary* > 1900") and act based on this



Transactions

