

Concurrency Control



Outline

- Concurrency control
 - Lock-Based Protocols (Two-Phase Locking)
 - Timestamp-Based Protocols
 - Multiversion Protocols (Snapshot Isolation)

Concurrency Control

- **Concurrency-control schemes** allow us to create, for a set of concurrent transactions, only acceptable schedules.
 - Isolation property may be weakened
 - They introduce an overhead

Tradeoff: amount of concurrency vs. amount of overhead

- **Goal:** Ideally, to reach the largest concurrency level whereas ensuring that executed schedule is:

serializable + recoverable + cascadeless

- A variety of concurrency-control schemes exists

No one is clearly the best

- Most used ones: **two-phase locking** and **snapshot isolation**



Lock-Based Protocols

- Access to data items in a **mutually exclusive** way
 - For a transaction to access an item, it needs to hold a **lock** on it
 - **Idea**: Hold the lock *long* enough to ensure serializability, but *short* enough to enhance concurrency
- Transactions can hold 2 **types of lock** on data :
 - a) **Shared** (S): Data item can only be *read* (no writing allowed).
 - **Several transactions** can have an S-lock on the same data item
 - b) **Exclusive** (X): Data item can be both *read and written*.
 - **Only one transaction** can have an X-lock on a data item at the same time
- If a lock holds on a data point p , another transaction may be granted a lock on it if the requested lock type is **compatible**:

	S	X
S	true	false
X	false	false

Compatibility table

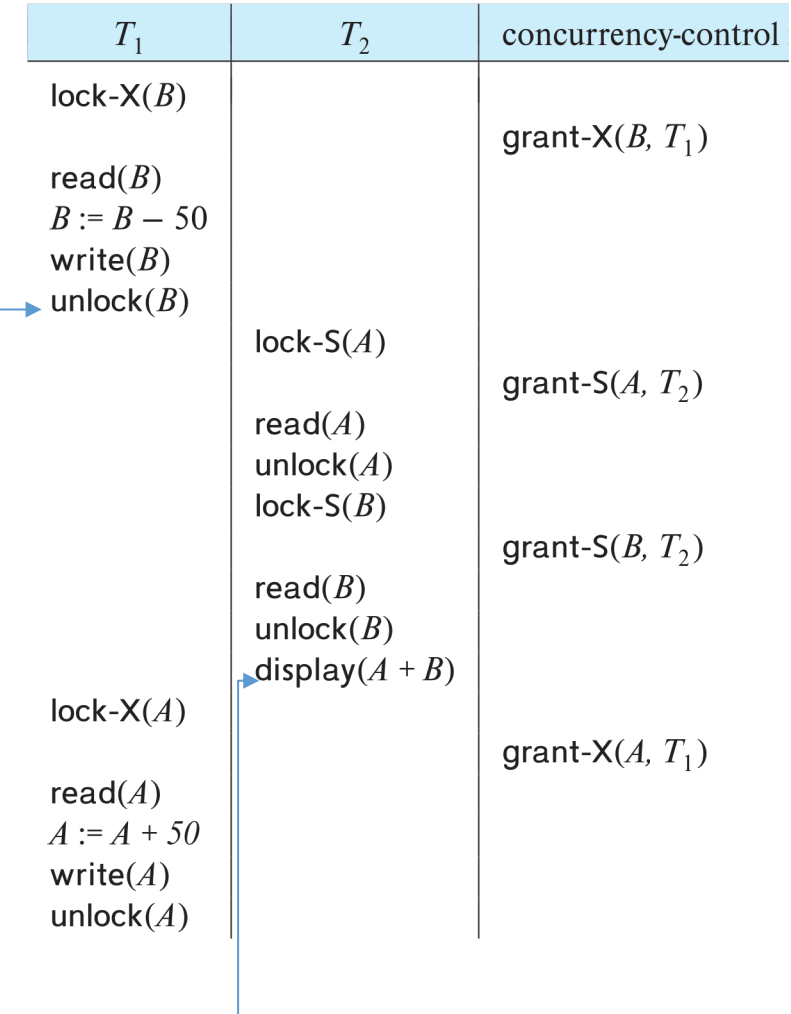


Lock-Based Protocols

- E.g., bank transaction:

T_1 : lock-X(B); read(B); $B := B - 50$; write(B); unlock(B); lock-X(A); read(A); $A := A + 50$; write(A); unlock(A);	T_2 : lock-S(A); read(A); unlock(A); lock-S(B); read(B); unlock(B); display(A+B);
--	--

- Simply locking is **not sufficient** to guarantee *serializability*
 - Unlock of B is too early
 - T_2 might display an inconsistent A+B value



Lock-Based Protocols

- E.g., bank transaction:

T_3 : lock-X(B); read(B); $B := B - 50$; write(B); lock-X(A); read(A); $A := A + 50$; write(A); unlock(B); unlock(A);	T_4 : lock-S(A); read(A); lock-S(B); read(B); display(A+B); unlock(A); unlock(B);
---	---

- Now, T_4 won't print inconsistent results!

Two-Phase Locking Protocol

- Protocol that ensures *conflict-serializable schedules*


- Two phases to issue lock/unlock requests:

1. Growing Phase

- Transaction can obtain locks
- Transaction **cannot** release locks

2. Shrinking Phase (after 1st lock release)

- Transaction **cannot** obtain locks
- Transaction can release locks



- **Lock point:** point where the last lock was acquired

- Transactions can be *ordered in terms of their lock points*
- This is a *serializability ordering* for the transactions!

Two-Phase Locking Protocol

- Susceptible of *deadlocks*.
 - **Deadlock**: no transaction can finish, as they mutually lock each other out
 - *Necessary evil*: deadlocks are preferable to inconsistent states
 - They can be handled by rolling back one of the transactions
- Susceptible of *starvation*.
- Susceptible of *cascading rollbacks*.

Deadlock	
T_3	T_4
lock-X(B) read(B) $B := B - 50$ write(B)	
	lock-S(A) read(A) lock-S(B)
lock-X(A)	

Two-Phase Locking Protocol

- Susceptible of *deadlocks*.
- Susceptible of *starvation*.
 - **Starvation**: when a transaction T_i requests an **X-lock** and needs to wait for a previous transaction's **S-lock** to be released, new **S-lock** requests of other transactions could advance T_i and block its access to the data
 - Possible solution: use **two conditions to grant a lock** on a data point p :
 1. There is no other transaction holding an *incompatible* lock on p
 2. There is no other transaction waiting for a lock on p
- Susceptible of *cascading rollbacks*.

Remember
compatibility table



Two-Phase Locking Protocol

- Susceptible of *deadlocks*.
- Susceptible of *starvation*.
- Susceptible of *cascading rollbacks*.

Possible solutions:

- **Strict two-phase locking:** a transaction must hold *all its exclusive locks* until it commits/aborts.
 - Ensures recoverability and avoids cascading rollbacks
- **Rigorous two-phase locking:** a transaction must hold *all locks* until it commits/aborts.
 - Transactions can be serialized in the order in which they commit.



Two-Phase Locking Protocol

- A transaction might need different lock types for the same data item in different moments.
- 2PL can sue **lock conversions**
 - **Upgrade** from a lock-S to a lock-X
 - Only during growing phase
 - **Downgrade** from a lock-X to a lock-S
 - Only during shrinking phase

Strict/rigorous two-phase locking, with **lock conversions**, are usually implemented in commercial DBMS

T_8	T_9
lock-S(a_1)	lock-S(a_1)
lock-S(a_2)	
lock-S(a_3)	
lock-S(a_4)	unlock(a_1) unlock(a_2)
lock-S(a_n) upgrade(a_1)	

Implementing 2PL

Automatic **generation of lock/unlock** requests

- When a transaction T_i makes a **read**(p),
 - it issues a **lock-S**(p) request
- When a transaction T_i makes a **write**(p),
 - if it holds a **lock-S**(p), it issues an **upgrade**(p) request;
 - Otherwise, it issues a **lock-X**(p) request
- All acquired locks are released when T_i commits/aborts

Implementing 2PL


- **Lock manager**: receives lock/unlock requests and answers with a:

A transaction waits until its request is answered

- Lock grant
- Order to roll back (deadlock)

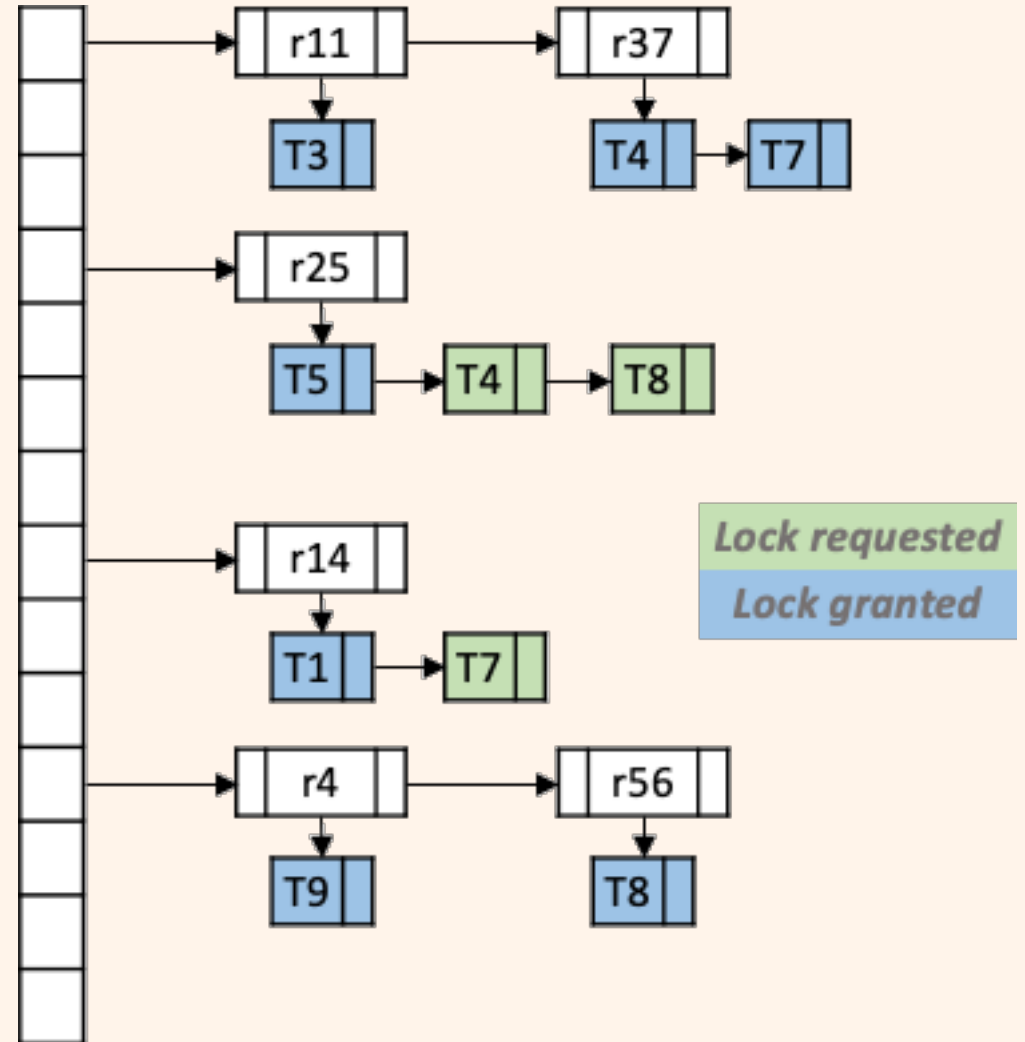
- Uses a **lock table**: hash structure to save granted locks and pending requests

- A queue for each data item (I_x)
- **Lock request** for I_x is appended to its queue of requests, and granted if it is compatible with any previous lock
- **Unlock request** leads to the deletion of the request, and checking if any request waits to be granted




Exercise

- What happens if T_9 requests a S-lock on item r_{37} ?
- Can T_7 request a X-lock on item r_4 ?
- What happens if T_3 requests a X-lock on item r_8 ?
- What happens if T_5 requests a X-lock on item r_{56} ?



Outline

- Concurrency control
 - Lock-Based Protocols (Two-Phase Locking)
 - Timestamp-Based Protocols 
 - Multiversion Protocols (Snapshot Isolation)

Timestamp-Based Protocols

- **Each transaction** T_i is given a timestamp $TS(T_i)$ as it enters the system.
 - Use system clock or a logical counter
- Manage concurrent execution such that
time-stamp order = serializability order
- **Each data point** p , is given **two** timestamps:
 - $W-TS(p)$: largest timestamp of any transaction that executed a **write**(p) successfully.
 - $R-TS(p)$: largest timestamp of any transaction that executed a **read**(p) successfully.

Timestamp-Ordering Protocol

- **Timestamp ordering (TSO)** protocol ensures that *conflicting operations* are executed in timestamp order.
- Transaction T_i can perform a **read**(p) only if no newer transaction has modified p
 - If $TS(T_i) < W-TS(p)$: read is *rejected*, and T_i is rolled back
 - If $TS(T_i) \geq W-TS(p)$: read is *executed*, and update
 $R-TS(p) := \max[R-TS(p) ; TS(T_i)]$
- T_i can perform a **write**(p) only if no newer transaction has manipulated (read or write) p
 - If $TS(T_i) < R-TS(p)$ or $TS(T_i) < W-TS(p)$: write is *rejected*, and T_i is rolled back
 - Otherwise, write is *executed*, and update
 $W-TS(p) := TS(T_i)$

Rolled back transactions are assigned new TSs




Exercise: TSO

- Is this schedule valid under TSO?
- Assume:
 - $R\text{-TS}(A) = W\text{-TS}(A) = 0$
 - $R\text{-TS}(B) = W\text{-TS}(B) = 0$
 - $\text{TS}(T_{25}) = 25$
 - $\text{TS}(T_{26}) = 26$

T_{25}	T_{26}
read(B)	read(B) $B := B - 50$ write(B)
read(A)	read(A)
display($A + B$)	$A := A + 50$ write(A) display($A + B$)

Timestamp-Ordering Protocol

- TSO protocol guarantees serializability since all the arcs in the precedence graph are of the form:



- Free of deadlocks, as no transaction ever waits.
- Schedules may **not** be *cascade-free* and even **not** *recoverable*.

Outline

- Concurrency control
 - Lock-Based Protocols (Two-Phase Locking)
 - Timestamp-Based Protocols
 - Multiversion Protocols (Snapshot Isolation)



Multiversion Protocols

- Issue: Read-only transactions that consult large amounts of data might conflict with update transactions

Poor performance

- Multiversion schemes maintain different **versions** of data points:
 - Each **write**(p) creates a new version of the data point, p .
 - A **read**(p) selects carefully the appropriate version of p , to ensure serializability.
 - Never waits, as the appropriate version is returned immediately.
- Several **variants**:
 - Multiversion Two-Phase Locking
 - Multiversion Timestamp Ordering
 - **Snapshot isolation**



Snapshot Isolation

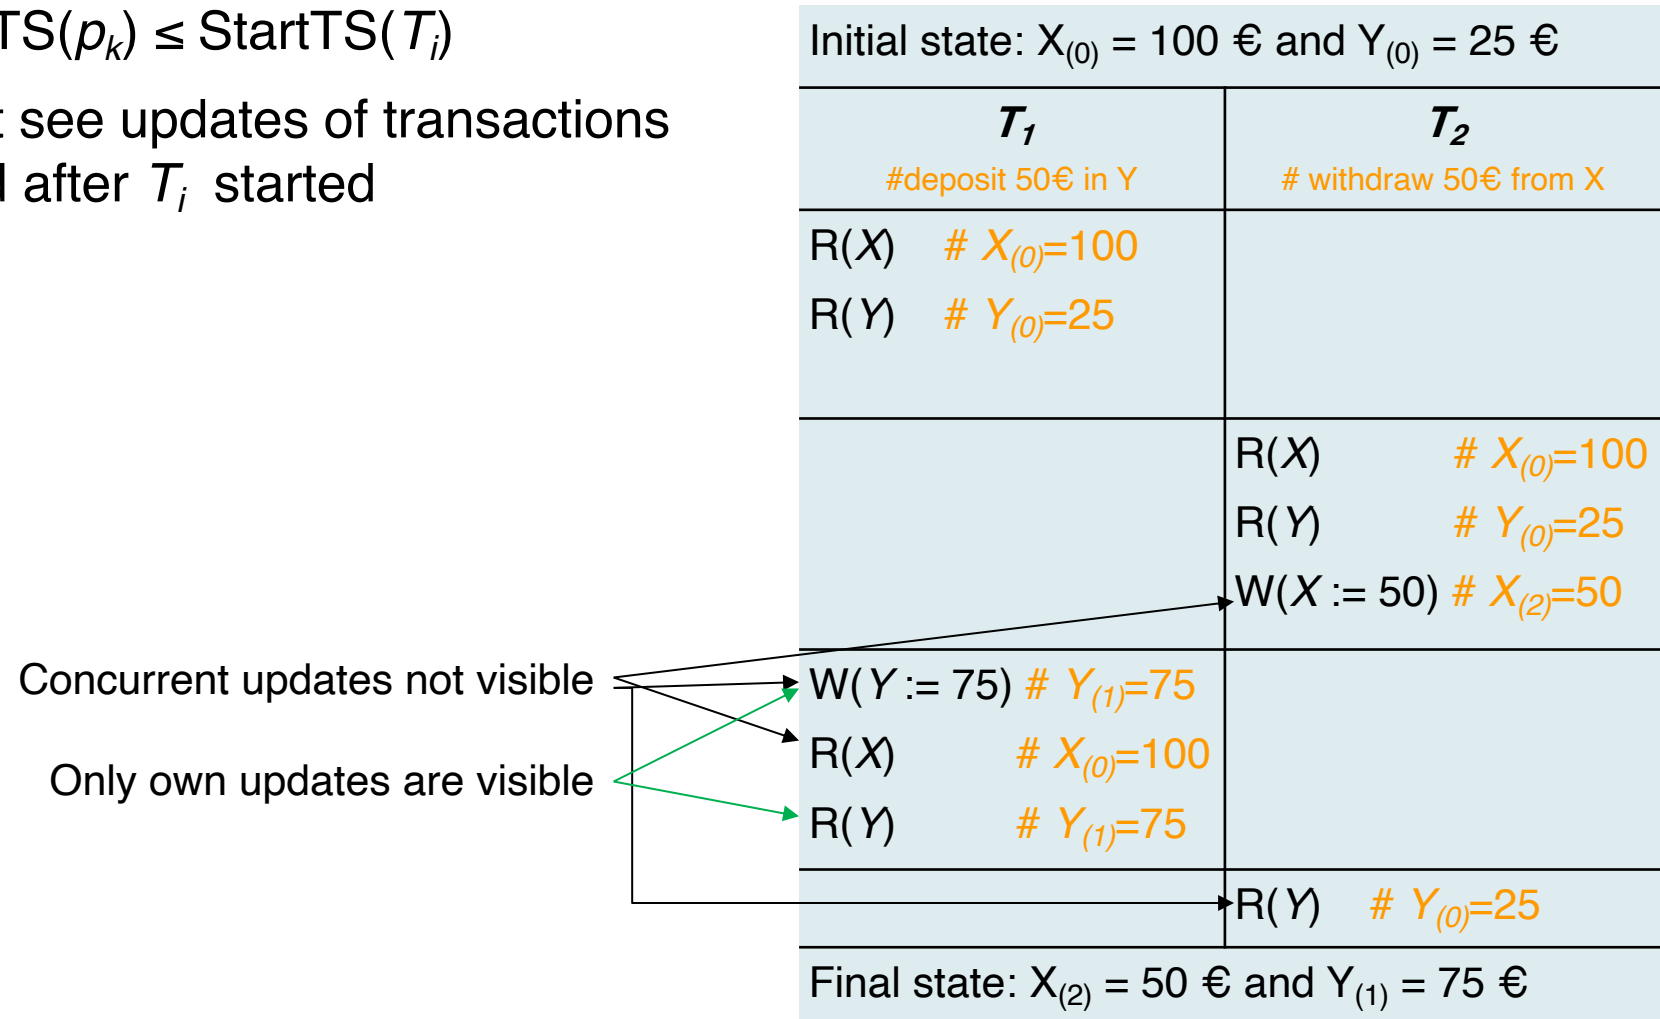
Widely used
in DBMS
(commercial/
open-source)

- Each transaction receives a “**snapshot**” (version) of the DB when it begins and operates locally
 - Transactions which modify the data need to be validated
 - After validation, allowed to commit and update the DB
 - Writing in the DB after commit needs to be an *atomic action*
- It uses timestamps too:
 - Each transaction T_i :
 - **StartTS**(T_i): when T_i started
 - **CommitTS**(T_i): when T_i requests validation
 - Each snapshot p_k for T_k :
 - **TS**(p_k) = **CommitTS**(T_k)



Snapshot Isolation

- When a transaction T_i reads a data point p , it uses the latest version p_k such that $TS(p_k) \leq \text{StartTS}(T_i)$
 - T_i doesn't see updates of transactions committed after T_i started



Snapshot Isolation

- T_i and T_j are **concurrent transactions** if:
 - $\text{StartTS}(T_j) \leq \text{StartTS}(T_i) \leq \text{CommitTS}(T_j)$, **or**
 - $\text{StartTS}(T_i) \leq \text{StartTS}(T_j) \leq \text{CommitTS}(T_i)$
- **Validating** a transaction for committing needs care:
 - 2+ *concurrent transactions* might have modified the same data point p .
 - Which one persists? Two strategies:
 - **First committer wins**
 - First updater wins


Snapshot Isolation

- **First-committer wins** rule:

T_i commits its copy p_i only if no other concurrent transaction T_k already wrote p with its own copy p_k

- Procedure

1. Check if data item p was written by a concurrent transaction:
 $\exists k : \text{StartTS}(T_i) \leq \text{TS}(p_k) \leq \text{CommitTS}(T_i) ?$
2. If so, T_i aborts
3. If not, T_i commits and p_i writes to the DB



Exercise SI

- Identify the values read by each operation of all the transactions
- Which commit request will be granted? Is there any that will be refused (+aborted)?

Before...: X=0 , Y=0 , Z=0		
T_1	T_2	T_3
Start		
W(Y := 1)		
W(X := 2)		
	R(X)	
	R(Z)	
	W(Y := 2)	
		Start
		R(Y)
	W(Z := 3)	
		W(Y := 3)
		R(Y)
R(Y)		
		R(Z)
R(Z)		
		Commit-Req
	Commit-Req	
Commit-Req		

Snapshot Isolation

- Useful when large reads are common
 - Reads are *never* blocked, and don't interfere with updates
- Avoids:
 - *Dirty reads* (reads uncommitted data)
 - *Lost updates* (updates overwritten by another transaction)
 - *Non-repeatable reads* (two reads of T_i see different values)
- **SI does not ensure serializability!!**
 - E.g., a serializable schedule would end up with $A=B$.
 - SI swaps the values of A and B
 - Note the cycle in the precedence graph!
 - This problem is known as **write skew**
 - Extension: **Serializable snapshot isolation (SSI)**

Set A = 3 and B = 17

T_i	T_j
read(A) read(B)	read(A) read(B)
$A=B$	$B=A$
write(A)	write(B)

Concurrency Control



Deadlock Handling

- **Deadlock** state refers to a *set of transactions* in which all them are waiting for another transaction in the set.

- Solution: rolling back (partially) some of the transactions

T_3	T_4
lock-X(B) read(B) $B := B - 50$ write(B)	
	lock-S(A) read(A) lock-S(B)
lock-X(A)	

- Two main methods:

- **Prevention protocols** ensure that the system will never enter a deadlock state
 - Useful if deadlocks are usual
 - **Detection-recovery scheme**: the system is allowed to enter in a deadlock state, but it tries to recover from it
 - More efficient if deadlocks are unusual
 - Involves overheads (run-time cost and potential losses)

Deadlock Handling: Prevention

Deadlock prevention strategies:

- Place locks intelligently
 - Transactions lock all its data items before it begins
 - You always lock data items placed after other items previously locked (ordering)
 - Needs to know which data items are going to be used
- Rollback instead of wait if expecting a deadlock
- Lock timeouts

Deadlock Handling: Prevention

Deadlock prevention strategies:

- Place locks intelligently
- Rollback instead of wait if expecting a deadlock
 - When T_i requests a lock on p , already locked by T_j :
 - **wait-die**: T_i waits if it is older than T_j ; or rolls back otherwise
 - **wound-die**: T_i waits if it is younger than T_j ; or rolls back otherwise
 - Might cause many unnecessary rollbacks
- Lock timeouts

Deadlock Handling: Prevention


Deadlock prevention strategies:

- Place locks intelligently
- Rollback instead of wait if expecting a deadlock
- **Lock timeouts** (less common, although easy to implement)
 - A transaction requests a lock and waits for **at most** a specific time.
 - If lock is not granted within that time, transaction *rolls itself back*
 - Might end up in *starvation*
 - Hard to decide **time** for timing out




Deadlock Handling: Detection-Recovery

- Periodically check for deadlocks and, if there is any, try to recover
- Use a *wait-for graph* to describe deadlocks:
 - *Vertices*: transactions
 - *Edges* ($T_i \rightarrow T_j$): if T_i is waiting for a lock held by T_j
 - *Deadlock state*: iff there is a cycle in the graph



Wait-for graph



Deadlocked wait-for graph

- *Deadlock-detection*: look for cycles in the graph
 - Periodicity depends on deadlocks' frequency

Deadlock Handling: Detection-Recovery

- *Recovery*, in case of deadlock detected:
 - **Victim selection**: transaction(s) to roll back (the one that incurs in “minimum cost”)
 - **Extent**: how far to roll back victim transaction?
 - **Total rollback**: Abort the transaction and restart it.
 - **Partial rollback**: Roll back only as far as necessary to break deadlock.
 - Release locks that other transactions are waiting for.
 - *Starvation* can happen if the victim is always the same!
 - Sol. 1: never select oldest transaction as the victim
 - Sol. 2: include no. rollbacks in cost function (*most common*)

