


Beyond traditional RDBMS



Historic view

- As of the rise of the Web 2.0:
 - Very large volumes of data being collected
 - Web logs first, then social media, web apps, etc.
 - Data from all type of sensors (phone, cars, ...)
 - Metadata from communication networks
 - Analytics on this data, of great value



- *Big Data*: different from earlier DBs by...
 - **Volume**: much larger amounts of data stored
 - **Velocity**: much higher data ingest rate
 - **Variety**: many data types, beyond relational data

Alternative DB systems

- Many apps. willing to sacrifice standard DB features (e.g., ACID properties) to enhance scalability
- Usually, an alternative is required when you need...
 - very high scalability
E.g., feed of a social network
 - to support non-relational data
E.g., video, documents, etc.
- Parallel processing is critical
- Different *alternative storage systems* include:
 - Sharding across multiple databases
 - Parallel and distributed databases
 - NoSQL models
 - Distributed file systems

Today

Rest of the semester

November 23-30



Sharding

- **Sharding**: partition data across **multiple databases**
- Split the DB on some **partitioning attributes**
 - *Range*: e.g., records with key in 1 to 100,000 on DB1, records with key in 100,001 to 200,000 on DB2, ...
 - *Hashing*
- **Properties:**
 - Scales well
 - Easy to implement
- **Drawbacks:**
 - Not transparent: application has to deal with routing of queries, queries spanning multiple DBs, etc.
 - If DB is overloaded, moving part of its load out is not straightforward nor easy
 - Chance of failure increases as more DBs are used
 - Use replicas (more work for application code)



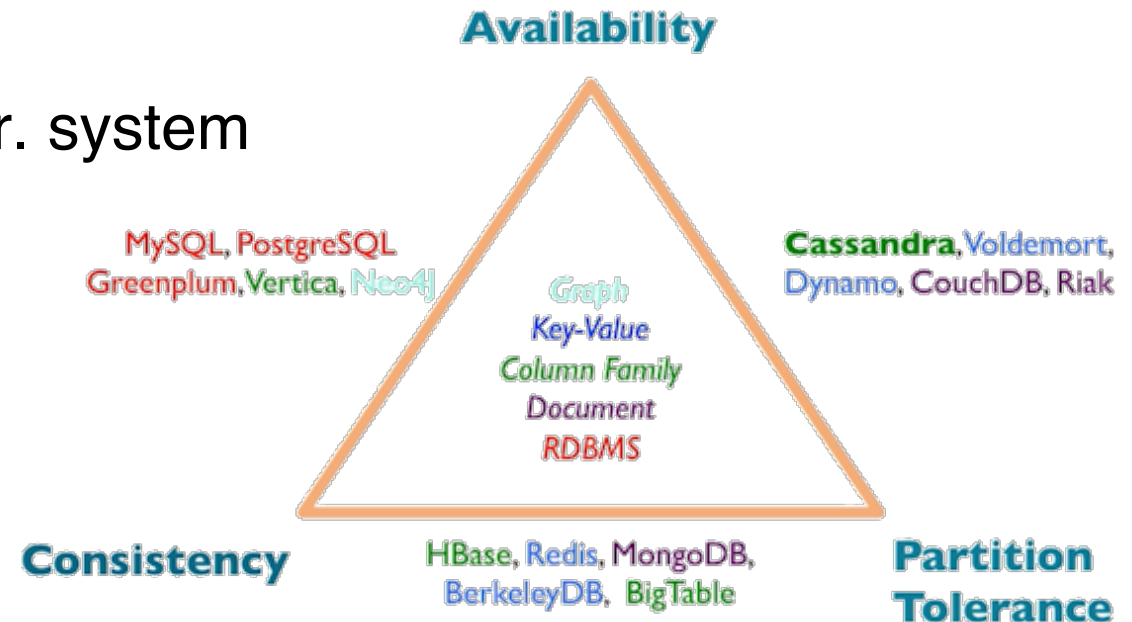
Availability vs. Consistency

- **Replication** is key for **availability** and essential in parallel/distributed DBs
- But, we would like to maintain replicated data **consistent**. Issues:
 - Atomic execution of update statement?
 - What if one of the systems that stores a copy is down?
 - You can use 3 replicas & read/write with 2+ replicas

- **Network partitions**: 2+ parts of a distr. system cannot communicate with each other

- In presence of **partitions**, **no protocol ensure both availability and consistency**
 - You need to decide which to ensure

CAP theorem



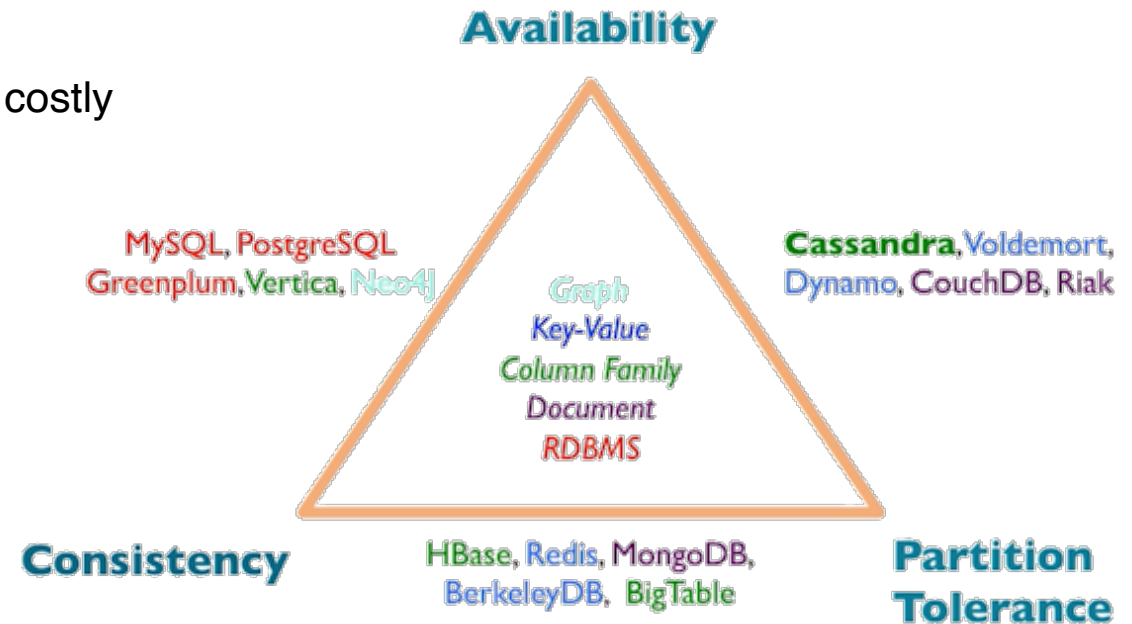
Availability vs. Consistency

CAP theorem

■ Network partitions

- In presence of **partitions**, no protocol ensure *both* **availability** and **consistency**
- In practice, in large systems **partitions** cannot be prevented.
 - You need to decide which to ensure: *consistency* or *availability*
 - Traditional DBMSs choose **consistency**
 - For distributed DBMSs, ensuring **consistency** is costly
 - Many applications, willing to reach higher **availability** at the cost of consistency

Really, **not a hard decision**
Consider *latency* instead of *availability*



Replication with Weak Consistency

- Key issues:
 - Some replicas may not get updated
 - Reads may get old versions
 - Different updates may be applied to different replicas
 - How to detect? Which is the newest version? How to resolve?

- Systems known as **BASE** (as opposed to ACID)

Basically Available, Soft state, Eventual consistency

- *Soft state*: copies of a data item may be inconsistent
- *Eventually consistent*: copies allowed to become inconsistent, but after partitioning is resolved, eventually all copies become consistent



Beyond traditional RDBMS

